

Report File

CSL4020

Assignment -1

Name: Ritam Sharma

(B20BB0030)

D=05,M=04,Y=02

A=0,B=3,C=0

ABC=030

FIRST=RITAM

LAST= SHARMA

PROG=BTECH

DATASET: FASHIONMNIST

WEIGHT INITIAL: Xavier(M=04)

data aug: rotate by 10 degrees and gaussian noise. (D=05)

Pooling: Avg pooling (M=04)

Target classification : (D+M+Y=5+4+2=11) your target 5 classes will be 1,3,5,7,9.

Question 1

Problem Statement

The goal is to train a CNN to classify 1,3,5,7,9. The data is augmented with GaussianNoise and random flipping

Hyperparameter tuning

1. Dataset consideration

The CNN model was trained firstly without augmented data and then secondly with the augmented data added, every other parameter is kept the same


Hyperparameters:

epoch	5
batch	100
optimizer	adam
learning rate	.001

▼ run 1

Architecture


Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 10, 26, 26]	100
BatchNorm2d-2	[-1, 10, 26, 26]	20
ReLU-3	[-1, 10, 26, 26]	0
MaxPool2d-4	[-1, 10, 13, 13]	0
Conv2d-5	[-1, 20, 11, 11]	1,820
BatchNorm2d-6	[-1, 20, 11, 11]	40
ReLU-7	[-1, 20, 11, 11]	0
Conv2d-8	[-1, 40, 9, 9]	7,240
BatchNorm2d-9	[-1, 40, 9, 9]	80
ReLU-10	[-1, 40, 9, 9]	0
Conv2d-11	[-1, 80, 7, 7]	28,880
BatchNorm2d-12	[-1, 80, 7, 7]	160
ReLU-13	[-1, 80, 7, 7]	0
MaxPool2d-14	[-1, 80, 3, 3]	0
Linear-15	[-1, 265]	191,065
Dropout2d-16	[-1, 265]	0
Linear-17	[-1, 5]	1,330
Total params: 230,735		
Trainable params: 230,735		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.40		
Params size (MB): 0.88		
Estimated Total Size (MB): 1.28		



Classification report for CNN :

	precision	recall	f1-score	support
0	0.99	0.98	0.98	270000
1	0.98	0.98	0.98	270000
2	0.97	0.97	0.97	270000
3	0.94	0.95	0.94	270000
4	0.96	0.95	0.96	270000
accuracy			0.97	1350000
macro avg	0.97	0.97	0.97	1350000
weighted avg	0.97	0.97	0.97	1350000

Accuracy of Trouser: 98.50%
Accuracy of dress: 98.93%
Accuracy of sandle: 97.03%
Accuracy of sneaker: 97.73%
Accuracy of ankle boot: 94.63%





▼ run 2

without augmented data same as run 1 (accuracy decreases by 1 %)

Classification report for CNN :

	precision	recall	f1-score	support
0	0.98	0.97	0.98	30000
1	0.97	0.98	0.98	30000
2	0.96	0.96	0.96	30000
3	0.93	0.94	0.94	30000
4	0.95	0.95	0.95	30000
accuracy			0.96	150000
macro avg	0.96	0.96	0.96	150000
weighted avg	0.96	0.96	0.96	150000





The addition of augmented data has increased the accuracy by 1 %, therefore for further analysis augmented data is added



2. epoch tuning

The CNN model was trained firstly with epochs 3,5,8,10 every other parameter is kept the same


Hyperparameters:

batch	100
optimizer	adam
learning rate	.001

▼ run 1 for epoch = 3





Classification report for CNN :				
	precision	recall	f1-score	support
0	0.98	0.97	0.98	30000
1	0.97	0.98	0.98	30000
2	0.96	0.96	0.96	30000
3	0.93	0.94	0.94	30000
4	0.95	0.95	0.95	30000
accuracy			0.96	150000
macro avg	0.96	0.96	0.96	150000
weighted avg	0.96	0.96	0.96	150000



▼ run 2 epoch=5

Architecture






Classification report for CNN :

	precision	recall	f1-score	support
0	0.99	0.98	0.98	270000
1	0.98	0.98	0.98	270000
2	0.97	0.97	0.97	270000
3	0.94	0.95	0.94	270000
4	0.96	0.95	0.96	270000
accuracy			0.97	1350000
macro avg	0.97	0.97	0.97	1350000
weighted avg	0.97	0.97	0.97	1350000

Accuracy of Trouser: 98.50%
 Accuracy of dress: 98.93%
 Accuracy of sandle: 97.03%
 Accuracy of sneaker: 97.73%
 Accuracy of ankle boot: 94.63%

▼ run 4 epoch=8



```
Accuracy of Trouser: 98.00%
Accuracy of dress: 99.57%
Accuracy of sandle: 97.27%
Accuracy of sneaker: 95.47%
Accuracy of ankle boot: 96.10%
```


```
Classification report for CNN :
              precision    recall  f1-score   support

0             0.99         0.98         0.98        432000
1             0.98         0.99         0.98        432000
2             0.97         0.96         0.97        432000
3             0.94         0.95         0.95        432000
4             0.96         0.95         0.96        432000

 accuracy                   0.97        2160000
 macro avg                  0.97         0.97        2160000
 weighted avg               0.97         0.97         0.97        2160000
```


▼ run 5 epoch=10

Very noisy learning curves



Classification report for CNN :				
	precision	recall	f1-score	support
0	0.99	0.98	0.99	540000
1	0.98	0.99	0.99	540000
2	0.98	0.97	0.97	540000
3	0.95	0.96	0.95	540000
4	0.96	0.96	0.96	540000
accuracy			0.97	2700000
macro avg	0.97	0.97	0.97	2700000
weighted avg	0.97	0.97	0.97	2700000

Accuracy of Trouser: 98.63%
 Accuracy of dress: 97.40%
 Accuracy of sandle: 95.43%
 Accuracy of sneaker: 96.30%
 Accuracy of ankle boot: 96.17%



The Best epoch was 5 , since for 10 the learning curves were really noisy but for 3 the accuracy was low

3. **Optimizer tuning:** The CNN model was trained firstly with SGD then Adam every other parameter is kept the same

▼ run 1 (SGD)

epoch: 5

batch : 100

optimizer : sgd


1 % dec than adam


learning rate : .001

Classification report for CNN :

	precision	recall	f1-score	support
0	0.99	0.97	0.98	270000
1	0.97	0.98	0.98	270000
2	0.96	0.95	0.95	270000
3	0.93	0.94	0.93	270000
4	0.95	0.95	0.95	270000
accuracy			0.96	1350000
macro avg	0.96	0.96	0.96	1350000
weighted avg	0.96	0.96	0.96	1350000

Accuracy of Trouser: 97.70%
Accuracy of dress: 98.63%
Accuracy of sandle: 96.10%
Accuracy of sneaker: 95.47%
Accuracy of ankle boot: 95.87%






▼ run 2 (Adam)


epoch: 5

batch : 100

optimizer : adam

learning rate : .001






Classification report for CNN :

	precision	recall	f1-score	support
0	0.99	0.98	0.98	270000
1	0.98	0.98	0.98	270000
2	0.97	0.97	0.97	270000
3	0.94	0.95	0.94	270000
4	0.96	0.95	0.96	270000
accuracy			0.97	1350000
macro avg	0.97	0.97	0.97	1350000
weighted avg	0.97	0.97	0.97	1350000

Accuracy of Trouser: 98.50%
 Accuracy of dress: 98.93%
 Accuracy of sandle: 97.03%
 Accuracy of sneaker: 97.73%
 Accuracy of ankle boot: 94.63%



Adam optimizer works better by 1% than SGD

Question 2

Problem Statement

The goal is to train an autoencoder and use the features from the encoder part to classify 1,3,5,7,9

The autoencoder is firstly trained to remove noise from data, then the classifier layer along with the encoder layer are trained to classify the classes 1 3 7 9

Hyperparameter tuning

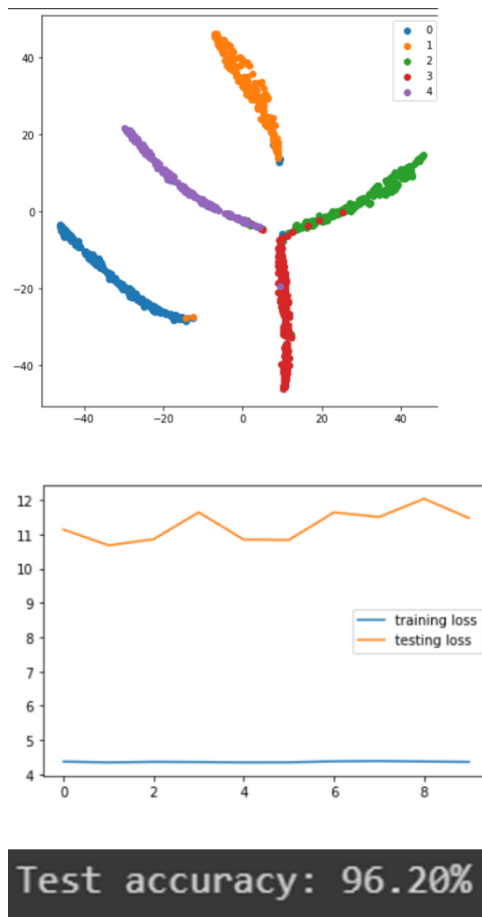
- learning rate tuning** For runs with the following hyperparameter, the following results was received

▼ run 1

Hyperparameters

learning rate	0.001
epoch	10



batch size	100
------------	-----



▼ run 2

Hyperparameters

learning rate	0.01
epoch	10
batch size	100



Test accuracy: 39.84%



Learning rate of 0.001 is preferred over 0.01 as the former has a much better tnse plot and accuracy

2. **model architecture tuning**, to models architecture, the following model architectures were tried

▼ run 1

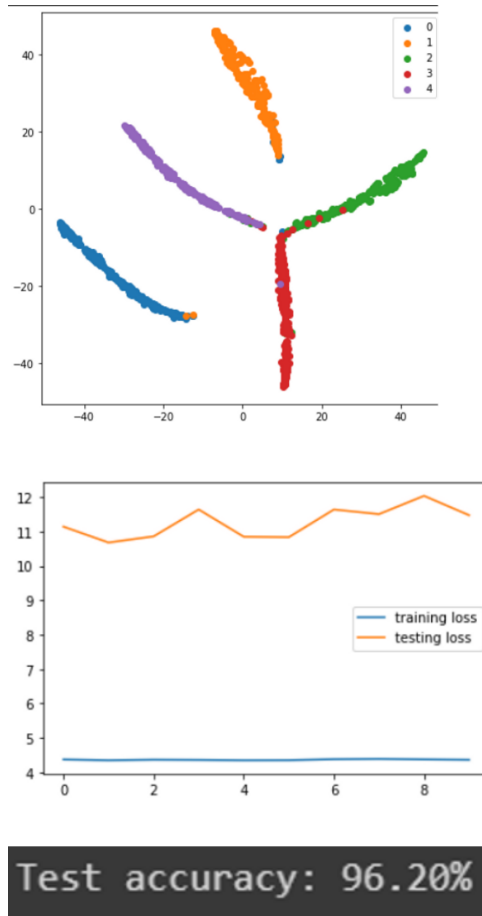
Architecture

of autoencoder since the classifier architecture is almost fixed

Layer (type)	Output Shape	Param #
Linear-1	[-1, 512]	401,920
ReLU-2	[-1, 512]	0
Linear-3	[-1, 256]	131,328
ReLU-4	[-1, 256]	0
Linear-5	[-1, 64]	16,448
ReLU-6	[-1, 64]	0
Linear-7	[-1, 256]	16,640
ReLU-8	[-1, 256]	0
Linear-9	[-1, 512]	131,584
ReLU-10	[-1, 512]	0
Linear-11	[-1, 784]	402,192
ReLU-12	[-1, 784]	0
Total params: 1,100,112		
Trainable params: 1,100,112		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.04		
Params size (MB): 4.20		
Estimated Total Size (MB): 4.24		

Hyperparameters

learning rate	0.001
epoch	10
batch size	100




▼ run 2

this is the thicker model

Layer (type)	Output Shape	Param #
Linear-1	[-1, 600]	471,000
ReLU-2	[-1, 600]	0
Linear-3	[-1, 550]	330,550
ReLU-4	[-1, 550]	0
Linear-5	[-1, 512]	282,112
ReLU-6	[-1, 512]	0
Linear-7	[-1, 550]	282,150
ReLU-8	[-1, 550]	0
Linear-9	[-1, 600]	330,600
ReLU-10	[-1, 600]	0
Linear-11	[-1, 784]	471,184
ReLU-12	[-1, 784]	0
Total params: 2,167,596		
Trainable params: 2,167,596		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.05		
Params size (MB): 8.27		
Estimated Total Size (MB): 8.33		

learning rate	0.001
epoch	10
batch size	100





the learning curve in run 2 is showing signs of overfitting since the test loss keeps increasing while the training loss is really low, furthermore, there is a decrease in accuracy on increasing the parameters therefore run1 has the better architecture

3. epoch tuning

▼ run1


train epoch : 2

Layer (type)	Output Shape	Param #
Linear-1	[-1, 512]	401,920
ReLU-2	[-1, 512]	0
Linear-3	[-1, 256]	131,328
ReLU-4	[-1, 256]	0
Linear-5	[-1, 64]	16,448
ReLU-6	[-1, 64]	0
Linear-7	[-1, 256]	16,640
ReLU-8	[-1, 256]	0
Linear-9	[-1, 512]	131,584
ReLU-10	[-1, 512]	0
Linear-11	[-1, 784]	402,192
ReLU-12	[-1, 784]	0
Total params: 1,100,112		
Trainable params: 1,100,112		
Non-trainable params: 0		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.04		
Params size (MB): 4.20		
Estimated Total Size (MB): 4.24		

An underfit model can be identified from the learning curve of the training loss only.

It may show a flat line or noisy values of relatively high loss, indicating that the model was unable to learn the training dataset at all.



therefore the following graph suggests that the the model is underfitting so we increase the epochs



Test accuracy: 94.36%

▼ run2

epoch : 2




Test accuracy: 96.20%

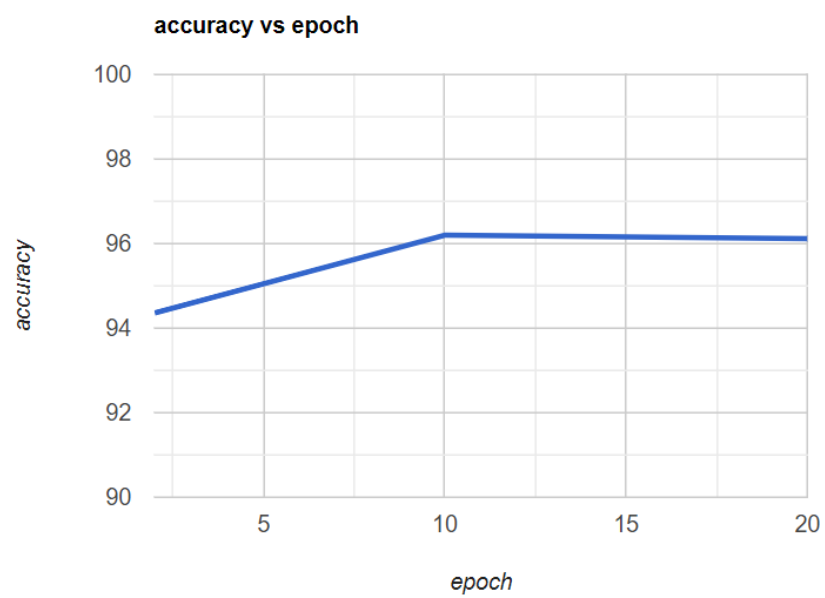

▼ run 3

epoch 20

(overfitting)



Test accuracy: 96.12%





as observed in the learning curves, epoch 2 was underfitting while epoch 20 was overfitting therefore 10 is the optimum epoch

Conclusion

CNN was performing quite better as compared to auto encoder