

LAB ASSIGNMENT 7

B20BB030

QUESTION OVERVIEW

The main goal is to Train ResNet-18 on the FashionMNIST dataset and Use Slurm to submit a job for training the model on the GPU server [15]

METRICS

ACCURACY

Accuracy is a useful metric because it provides a straightforward measure of how well a model is performing on a given task. However, it is important to note that accuracy may not be the best metric in all cases, as it does not take into account the relative costs of different types of errors.

For example, in the context of Fashion-MNIST, misclassifying a sandal as a sneaker may not be as costly as misclassifying a coat as a T-shirt. Other metrics such as precision, recall, and F1 score may be more appropriate in such cases.

Overall, accuracy is a useful metric for evaluating the performance of machine learning models on the Fashion-MNIST dataset, but it should be used in conjunction with other metrics to provide a more complete picture of the model's performance.

METHODOLOGY

TO TRAIN FASHION MNIST ON RESNET

DATA PREPROCESSING:

First, you need to preprocess the Fashion-MNIST dataset. This includes resizing the images to a uniform size (e.g., 224x224), normalizing the pixel values to be between 0 and 1, and dividing the dataset into training and validation sets.

MODEL ARCHITECTURE:

The ResNet model architecture is a deep convolutional neural network that has been proven effective for image classification tasks. You can use an existing ResNet architecture, such as ResNet-50 or ResNet-101, or you can create a custom ResNet architecture.

INITIALIZATION:

You can initialize the weights of the ResNet model using a pre-trained model, such as ImageNet, to speed up the training process and improve the performance of the model.

TRAINING:

You can train the ResNet model on the Fashion-MNIST dataset using a suitable optimizer such as stochastic gradient descent (SGD) or Adam. You will also need to define the loss function, such as categorical cross-entropy, and choose an appropriate learning rate for the optimizer.

EVALUATION:

You can evaluate the performance of the trained ResNet model on the validation set using metrics such as accuracy, precision, recall, and F1 score. You can also visualize the performance of the model using plots of the loss and accuracy over time.

HYPERPARAMETER TUNING:

You can fine-tune the hyperparameters of the ResNet model, such as the learning rate, batch size, and a number of epochs, to improve its performance on the validation set. This can be done using techniques such as grid search or random search.

TESTING:

Finally, you can test the performance of the trained ResNet model on a separate test set to get a final measure of its performance. This can be done by loading the saved model weights, running the model on the test set, and then evaluating its performance using metrics such as accuracy.

TO USE SLURM TO SUBMIT A JOB FOR TRAINING THE MODEL ON THE GPU SERVER

COPY FASHION MNIST DATA FROM LOCAL TO GPU SERVER USING SCP -R COMMAND

```
C:\Users\Ritam>scp -r C:\Users\Ritam\OneDrive\Documents\delete\archive b20bb030@172.25.0.15:~/scratch/data/b20bb030/lab7
b20bb030@172.25.0.15's password:
archive.zip                                100% 59MB 105.7MB/s 00:00
fashion-mnist_test.csv                    100% 21MB 111.7MB/s 00:00
fashion-mnist_train.csv                   100% 827MB 108.8MB/s 00:01
t10k-images-idx3-ubyte                    100% 7656KB 88.8MB/s 00:00
t10k-labels-idx1-ubyte                    100% 10KB 9.8KB/s 00:00
train-images-idx3-ubyte                    100% 45MB 106.9MB/s 00:00
train-labels-idx1-ubyte                    100% 59KB 58.6KB/s 00:00
C:\Users\Ritam>
```

CONNECT TO THE HOST ON VSCODE BY TYPING THE PASSWORD

```
main.py - lab7 [SSH: 172.25.0.15] - Visual Studio Code
File Edit Selection View Go Run Terminal Help
EXPLORER
  LAB7 [SSH: 172.25.0.15]
    > Accuracy plot
    > archive
    > loss plot
    main.py
    my_list.pickle
    result_40460.log
    result.log
    slurm.sh
    test.log
    test.py
main.py
49
50 # Define the ResNet-18 model
51 model = torch.hub.load('pytorch/vision:v0.9.0', 'resnet18', pretrained=False, num_classes=1000)
52 model=model.to(device)
53
54 # Define the loss function and optimizer
55 criterion = nn.CrossEntropyLoss()
56
57 # Train the model
58 # Set hyperparameters
59 num_epochs_list = [5, 10, 15]
60 learning_rate_list = [0.001, 0.01, 0.1]
61 batch_size_list = [64, 128, 256]
62 result=[]
63 for num_epochs in num_epochs_list:
64
65     for learning_rate in learning_rate_list:
66         for batch_size in batch_size_list:
67             # Train the model
68             optimizer = optim.Adam(model.parameters(), lr=learning_rate)
69
```

GENERATE A SLURM.SH FILE AND A MAIN SCRIPT AS SHOWN

```
slurm.sh
1  #!/bin/bash
2  #SBATCH --job-name=test_job      # Job name
3  #SBATCH --partition=gpu2        #Partition name can be test/small/medium/large/gpu #Partition
4  #SBATCH --nodes=1              # Run all processes on a single node
5  #SBATCH --ntasks=1             # Run a single task
6  #SBATCH --cpus-per-task=4       # Number of CPU cores per task
7  #SBATCH --gres=gpu             # Include gpu for the task (only for GPU jobs)
8  #SBATCH --mem=6gb              # Total memory limit
9  #SBATCH --time=10:05:00        # Time limit hrs:min:sec
10 #SBATCH --output=result_%j.log  # Standard output and error log
11
12 module load python/3.8
13
14 module load openmpi4
15 nvidia-smi
16 #cd /iitjhome/b20bb030/scratch/data/b20bb013
17 python3.8 main.py
```


SUBMIT THE BATCH JOB BY SBATCH SLURM.SH COMMAND AND OPEN THE RESULT.LOG

```
result_40460.log - lab7 [SSH: 172.25.0.15] - Visual Studio Code
File Edit Selection View Go Run Terminal Help
EXPLORER
  LAB7 [S...
    > Accuracy plot
    > archive
    > loss plot
    main.py
    my_list.pickle
    result_40460.log
    result.log
    slurm.sh
    test.log
    test.py
main.py
result_40460.log
20 GPU GI CI PID Type Process name GPU Memory
21 ID ID ID Type Process name Usage
22 =====
23 | No running processes found
24 +-----+
25 Using cache found in /iitjhome/b20bb030/.cache/torch/hub/pytorch_vision_v0.9.0
26 /scratch/apps/python/3.8/lib/python3.8/site-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pr
27 warnings.warn(
28 /scratch/apps/python/3.8/lib/python3.8/site-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other t
29 warnings.warn(msg)
30 File running
31 Epoch [1/5], Loss: 0.5291, Accuracy: 78.98%
32 Epoch [2/5], Loss: 0.4822, Accuracy: 82.83%
33 Epoch [3/5], Loss: 0.4318, Accuracy: 84.03%
34 Epoch [4/5], Loss: 0.3837, Accuracy: 85.78%
35 Epoch [5/5], Loss: 0.4038, Accuracy: 85.64%
36 105.67938342876732
PROBLEMS OUTPUT TERMINAL PORTS DEBUG CONSOLE
[lab7] bash - lab7
• [b20bb030@hpclogin lab7]$ sbatch slurm.sh
Submitted batch job 40455
• [b20bb030@hpclogin lab7]$ scancel 40455
• [b20bb030@hpclogin lab7]$ sbatch slurm.sh
Submitted batch job 40457
• [b20bb030@hpclogin lab7]$ sbatch slurm.sh
Submitted batch job 40460
• [b20bb030@hpclogin lab7]$
```


RESULTS

In this Section, I have plotted all the loss curves and accuracy curves for the varying hyperparameters along with a table that summarizes the findings


LOSS CURVES [EPOCH NUMBER, LEARNING RATE, BATCH SIZE, ACCURACY, TIME ELAPSED IN CAPTION]




loss curve for 10 , 0.001, 128,
85.56666666666666,184.973436685279




loss curve for 10 , 0.001, 256,
85.95166666666667,171.61322858929634




loss curve for 10 , 0.001, 64,
84.23833333333333,214.85168339684606




loss curve for 10 , 0.01, 128,
87.43333333333334,183.06576551962644




loss curve for 10 , 0.01, 256, 87.41,171.71338925231248




loss curve for 10 , 0.01, 64, 86.9,208.79939623270184




loss curve for 10 , 0.1, 128, 85.66666666666667,189.7662300132215




loss curve for 10 , 0.1, 256, 87.34666666666666,175.1224684547633




loss curve for 10 , 0.1, 64, 85.5016666666667,215.26309788413346




loss curve for 15 , 0.001, 128,
88.05166666666666,266.960066286847




loss curve for 15 , 0.001, 256,
88.16166666666666,250.16800520848483




loss curve for 15 , 0.001, 64,
87.68666666666667,302.57188345771283



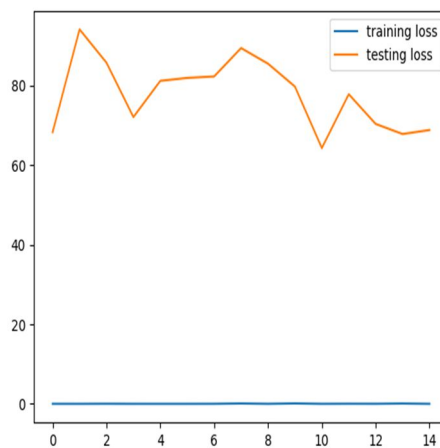
loss curve for 15 , 0.01, 128,
87.88333333333334,271.27991808764637



loss curve for 15 , 0.01, 256, 87.94,251.83222610410303



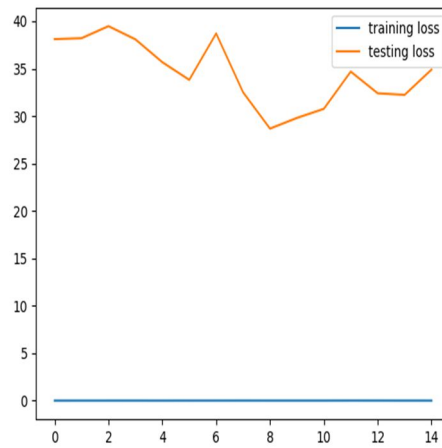
loss curve for 15 , 0.01, 64, 87.72333333333333,301.6737953508273



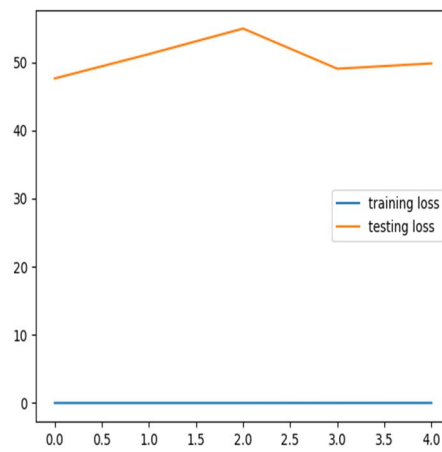
loss curve for 15 , 0.1, 128, 87.095,268.8499610843137



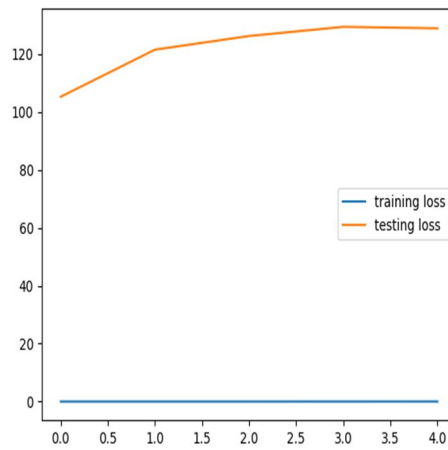
loss curve for 15 , 0.1, 256, 87.305,250.71896273549646



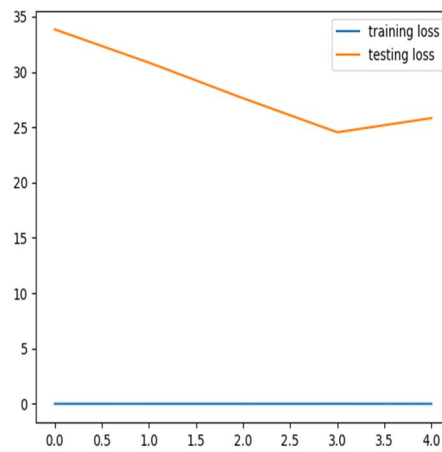
loss curve for 15 , 0.1, 64, 86.57833333333333,303.23279016278684




loss curve for 5 , 0.001, 128, 87.54,94.45059104356915



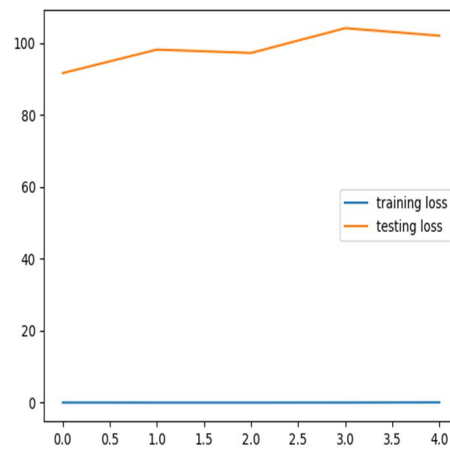
loss curve for 5, 0.001, 256, 87.45833333333333,85.90150786098093




loss curve for 5, 0.001, 64, 85.63833333333334,105.67938342876732




loss curve for 5, 0.01, 128, 86.265,91.20869947131723



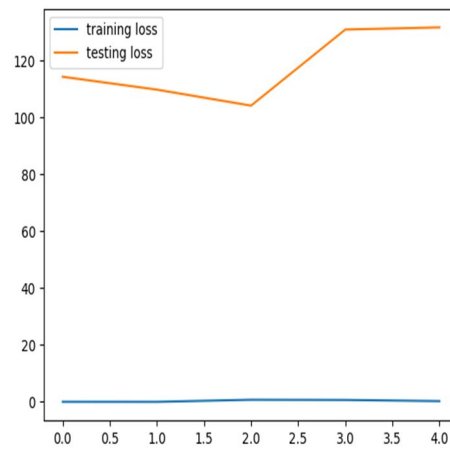
loss curve for 5, 0.01, 256, 86.90333333333334,85.85660909023136



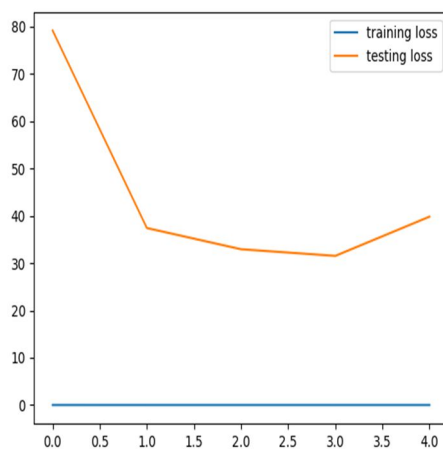
loss curve for 5, 0.01, 64, 83.21666666666667, 106.19132073782384



loss curve for 5, 0.1, 128, 83.93333333333334, 91.82626716885716

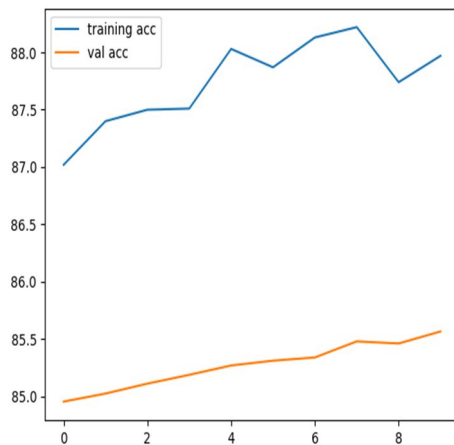


loss curve for 5, 0.1, 256, 81.855,85.9542280882597

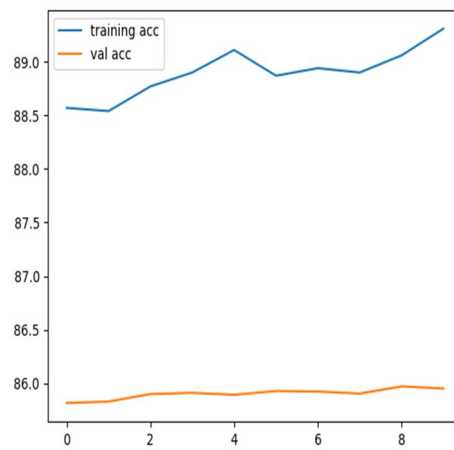


loss curve for 5, 0.1, 64, 78.30333333333333,107.80648355931044


ACCURACY CURVES [EPOCH NUMBER, LEARNING RATE, BATCH SIZE, ACCURACY, TIME ELAPSED
IN CAPTION]



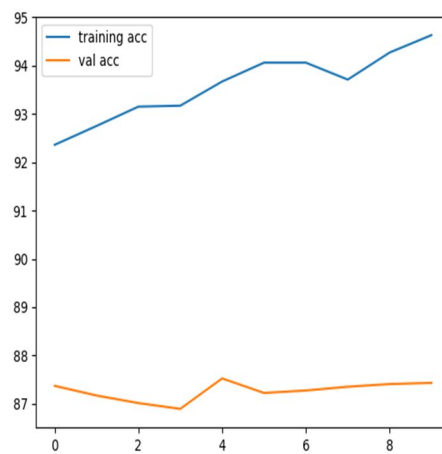
ACCURACY curve for 10, 0.001, 128,85.56666666666666,184.973436685279




ACCURACY curve for 10, 0.001, 256, 85.95166666666667, 171.61322858929634




ACCURACY curve for 10, 0.001, 64, 84.23833333333333, 214.85168339684606




ACCURACY curve for 10, 0.01, 128, 87.43333333333334, 183.06576551962644




ACCURACY curve for 10 , 0.01, 256,87.41.171.71338925231248




ACCURACY curve for 10 , 0.01, 64,86.9.208.79939623270184




ACCURACY curve for 10 , 0.1, 128,85.66666666666667.189.7662300132215




ACCURACY curve for 10 , 0.1, 256,87.34666666666666.175.1224684547633




ACCURACY curve for 10 , 0.1, 64,85.50166666666667.215.26309788413346




ACCURACY curve for 15 , 0.001, 128,88.05166666666666.266.960066286847




ACCURACY curve for 15 , 0.001, 256,88.16166666666666.250.16800520848483




ACCURACY curve for 15 , 0.001, 64,87.68666666666667.302.57188345771283




ACCURACY curve for 15 , 0.01, 128,87.88333333333334.271.27991808764637




ACCURACY curve for 15 , 0.01, 256,87.94.251.83222610410303




ACCURACY curve for 15 , 0.01, 64,87.72333333333333.301.6737953508273




ACCURACY curve for 15 , 0.1, 128,87.095.268.8499610843137




ACCURACY curve for 15 , 0.1, 256,87.305.250.71896273549646




ACCURACY curve for 15 , 0.1, 64,86.57833333333333.303.23279016278684




ACCURACY curve for 5 , 0.001, 128,87.54.94.45059104356915




ACCURACY curve for 5 , 0.001, 256,87.45833333333333.85.90150786098093




ACCURACY curve for 5 , 0.001, 64,85.63833333333333.105.67938342876732




ACCURACY curve for 5 , 0.01, 128,86.265.91.20869947131723




ACCURACY curve for 5 , 0.01, 256,86.90333333333334.85.85660909023136




ACCURACY curve for 5 , 0.01, 64,83.21666666666667.106.19132073782384



ACCURACY curve for 5 , 0.1, 128,83.93333333333334.91.82626716885716



ACCURACY curve for 5 , 0.1, 256,81.855.85.9542280882597



ACCURACY curve for 5 , 0.1, 64,78.30333333333333.107.80648355931044

SUMMARY OF RESULTS IN TABULAR FORM

TIMEELAPSED(SECONDS)	EPOCH	LEARNING RATE	BATCH SIZE	ACCURACY
105.679383	5	0.001	64	85.638333
94.450591	5	0.001	128	87.540000
85.901508	5	0.001	256	87.458333
106.191321	5	0.010	64	83.216667

TIME ELAPSED(SECONDS)	EPOCH	LEARNING RATE	BATCH SIZE	ACCURACY
91.208699	5	0.010	128	86.265000
85.856609	5	0.010	256	86.903333
107.806484	5	0.100	64	78.303333
91.826267	5	0.100	128	83.933333
85.954228	5	0.100	256	81.855000
214.851683	10	0.001	64	84.238333
184.973437	10	0.001	128	85.566667
171.613229	10	0.001	256	85.951667
208.799396	10	0.010	64	86.900000
183.065766	10	0.010	128	87.433333
171.713389	10	0.010	256	87.410000
215.263098	10	0.100	64	85.501667
189.766230	10	0.100	128	85.666667
175.122468	10	0.100	256	87.346667
302.571883	15	0.001	64	87.686667
266.960066	15	0.001	128	88.051667
250.168005	15	0.001	256	88.161667
301.673795	15	0.010	64	87.723333
271.279918	15	0.010	128	87.883333
251.832226	15	0.010	256	87.940000

TIME ELAPSED (SECONDS)	EPOCH	LEARNING RATE	BATCH SIZE	ACCURACY
303.232790	15	0.100	64	86.578333
268.849961	15	0.100	128	87.095000
250.718963	15	0.100	256	87.305000

DISCUSSION OF THE RESULTS

- Based on the highlighted results for epoch 5, the model achieved an accuracy of around 87.5% (best) after 5 epochs with a learning rate of 0.001 and batch sizes of 128 and 256. The elapsed time for training the model with batch size 256 was lower than that of batch size 128, therefore a batch size of 256 seems to be optimal. Decreasing the batch size to 64 didn't improve the accuracy either.
- Similarly, in the case of 10 epochs, a batch size of 256 gave comparable results to that of 128 with only a slight decrease in accuracy (87.433333 vs 87.410000) but the time elapsed reduced by approximately 12 seconds.
- Meanwhile increasing the batch size from 128 to 256 for 15 epochs actually increased the accuracy (by 0.11%) and even decreased the time elapsed by 16 seconds !
- Therefore it is concluded from these runs, that a batch size of 256 is optimal for my model, and for the following epochs, the optimal hyperparameters are given below:

EPOCH	LEARNING RATE	BATCH SIZE
5	0.001	256
10	0.010	256
15	0.001	256